From Revenue Volatility to Resilience Engineering: A Behavioral Modeling Approach for Constrained Retail Optimization

Sirikonda Vinod, Shaik Abdul Razak
Team UVR TECHNIES
N.L. Dalmia Institute of Management Studies & Research
Mumbai, India
{vinod.s, abdul.r}@uvrtechnies.in

***Abstract*—The contemporary retail landscape faces unprecedented challenges from revenue concentration risk, AOV compression, and binding capital constraints. This paper presents a comprehensive behavioral revenue modeling framework applied to STABILIS, a multi-category retail platform experiencing 22.7% revenue contraction with only 7% user base decline. Through systematic exploratory data analysis linked to predictive modeling, we identify that 28% of users (Champions) drive 68.1% of total revenue, establishing concentration as the primary systemic vulnerability. We develop and validate an Expected Revenue Per User (ERPU) framework decomposing revenue into purchase probability, expected basket value, and return risk components. The model demonstrates robust generalization with 2.64x lift in top decile validation (improving from 2.1x training lift). Under a binding R 25 million capital constraint with 15% minimum margin threshold, greedy optimization achieves 116.8x ROI by targeting 114 top Champion users, generating R 2.92 billion projected revenue (64% growth). Critical sensitivity analysis reveals 5% Champion churn causes R 138 million loss, necessitating VIP retention protocols. This research contributes a replicable framework linking exploratory visualization to constrained optimization, with implications for retailers facing similar structural contractions.**

***Index Terms*—Behavioral Revenue Modeling, Customer Concentration Risk, Constrained Optimization, RFM Segmentation, ERPU Framework, Retail Analytics, Capital Allocation, Sensitivity Analysis**

I. INTRODUCTION

The rapid diffusion of digital technologies has fundamentally reshaped consumer behavior patterns while simultaneously exposing retailers to unprecedented revenue volatility [1]. Management faces a critical paradox: user bases remain relatively stable while revenue streams exhibit sharp contractions, indicating structural shifts in purchase behavior rather than simple churn dynamics [2].

STABILIS operates a multi-category retail platform where management has observed four critical phenomena: uneven purchase frequency across users, significant basket size variability, concentrated revenue among a small user segment, and return-prone behavior eroding margins [3]. These observations mirror broader industry challenges documented in omnichannel retail literature, where capacity constraints and return rates fundamentally alter inventory economics [5].

The research questions guiding this investigation are:

RQ1: How can exploratory data visualization systematically identify behavioral drivers of revenue concentration and AOV compression?

RQ2: What modeling framework effectively decomposes retail revenue into predictive components while maintaining interpretability for board-level decision-making?

RQ3: Under binding capital constraints, how should retailers optimize customer targeting to maximize contribution margin while managing sensitivity risks?

Our contributions are threefold. First, we demonstrate a systematic methodology linking exploratory graphics (histograms, box plots, mosaic plots, scatter matrices) to predictive model selection, establishing visual-to-analytical translation protocols. Second, we develop and validate an ERPU framework that separates probability modeling from value modeling, incorporating Laplace smoothing for frequency estimation, 95th percentile basket clipping for heavy-tail control, and return risk anchoring for zero-history users. Third, we implement constrained greedy optimization under a R 25 million capital constraint with 15% margin threshold, achieving 116.8x ROI while documenting critical sensitivity to Champion churn (R 138 million loss per 5% churn).

The paper is structured as follows: Section II reviews relevant literature on omnichannel retail operations, customer concentration, and constrained optimization. Section III presents our data structure and exploratory methodology. Section IV develops the behavioral revenue modeling framework. Section V details validation results and model recalibration. Section VI presents constrained optimization outcomes and sensitivity analysis. Section VII discusses managerial implications and limitations. Section VIII concludes with recommendations for retailers facing structural contractions.

II. LITERATURE REVIEW

A. Omnichannel Retail Operations and Revenue Volatility

The expansion of omnichannel retail practices has garnered substantial scholarly attention, particularly regarding how retailers should integratively manage online and physical channels to deliver seamless shopping experiences [1]. A key characteristic is consumers' ex-ante value uncertainty, as certain non-digital product attributes cannot be fully assessed online [6]. Offline consumers resolve this uncertainty through physical inspection, while online shoppers may require returns when actual value falls short of expectations [5].

Gallino and Moreno [1] empirically demonstrated cross-selling and channel-migration effects in omnichannel contexts. Gao and Su [6] theoretically analyzed flexible fulfillment policies, noting that inventory visibility might paradoxically reduce store visits during stockouts, diminishing cross-selling opportunities. Shi et al. [7] examined omnichannel implementations under both resalable and non-resalable return conditions, deriving optimal two-period pricing and inventory targets.

Our work extends this literature by explicitly modeling revenue concentration risk—the phenomenon where a minority of customers drive majority revenue—which prior studies have addressed only tangentially. Li et al. [3] examined consumer heterogeneity in e-retailing markets, finding that risk-taking tendencies influence market concentration dynamics. Similarly, Teng Li [3] demonstrated through agent-based simulation that consumer pickiness and word-of-mouth weight significantly affect concentration patterns, with negative ratings suppressing sales advantages.

B. Customer Concentration and Revenue Risk

Revenue concentration represents both an efficiency opportunity and a systemic vulnerability. The Pareto principle (80/20 rule) manifests strongly in retail contexts, but its implications for capital-constrained optimization remain underexplored. Recent industry research indicates that inventory distortion costs global retailers \$1.73 trillion annually, equivalent to 6.5% of worldwide retail sales [26]. Brands deploying AI-driven inventory management achieve 2.3x higher sales growth and 2.5x higher profit growth than competitors, yet less than 25% of retailers have successfully implemented such systems [27].

The Cyber 5 2025 shopping period provided critical insights: despite 15% traffic decline, ordered revenue climbed 12% for brands with synchronized advertising and inventory [28]. Revenue loss from out-of-stock items plummeted 54% year-over-year for aligned brands, while misaligned categories (furniture, grocery) experienced 269% and 123% OOS loss spikes respectively. These findings underscore that concentration risk management requires real-time visibility across advertising, inventory, and fulfillment systems.

C. Constrained Optimization in Customer Targeting

Modern e-commerce services frequently target customers with incentives to increase acquisition and retention [4]. However, random targeting or propensity-based approaches can be suboptimal, as they fail to identify customers who would benefit most from intervention while respecting operational constraints [25].

Li et al. [4] proposed a two-stage policy framework combining uplift modeling with constrained optimization, demonstrating improvement over state-of-the-art approaches through large-scale experimental studies. Their methodology first estimates conditional average treatment effects (CATE) using meta-learners or forest-based estimators [15], then solves:

$$\max_{\pi} \sum_{i} \sum_{k} \pi_{ik} w_i \hat{\tau}_k(x_i) \tag{1}$$

$$\text{subject to } \pi_{ik} \in \{0, 1\}, \sum_{k} \pi_{ik} = 1, g(\pi) \leq c \tag{2}$$

Our optimization framework builds directly on this approach, adapting it for retail revenue maximization under capital constraints. Unlike prior work focusing on binary treatment effects, we model continuous revenue outcomes through ERPU decomposition, enabling finer-grained resource allocation.

D. Behavioral Revenue Decomposition

The separation of purchase probability from basket value modeling represents a fundamental analytical discipline [19]. Probability modeling addresses customer engagement likelihood, while value modeling addresses transaction magnitude. Return risk introduces a third dimension, capturing margin erosion [21].

Laplace smoothing addresses sparse data challenges for low-frequency users, preventing overfitting to noise. Basket clipping at the 95th percentile controls heavy-tail distortion, where extreme spenders can skew predictions and lead to over-investment in volatile segments [22]. Return anchoring at global means provides conservative estimates for users without return history.

E. Research Gap

Despite substantial progress in omnichannel operations and constrained optimization, no unified framework exists linking exploratory visualization to behavioral revenue decomposition under capital constraints. Prior studies either focus on channel strategy [1], consumer heterogeneity [3], or optimization methodology [4] in isolation. This paper addresses this gap by demonstrating an end-to-end methodology applicable to retailers facing structural revenue contractions.

III. DATA AND EXPLORATORY METHODOLOGY

A. Dataset Description

The analysis utilizes two transactional datasets provided by STABILIS:

TABLE I
DATASET OVERVIEW

Dataset	Period	Transactions	Unique Customers	Unique Pr
Customer Transactions	2019-2020	417,534	4,383	2,031
Customer Validation	2020-2021	323,055	4,068	930 valid

TABLE II
VARIABLE STRUCTURE

Variable	Type	Level	Description
EventID	Nominal	Basket	Unique basket identifier
EventType	Binary	Line Item	Purchase (+) / Return (-)
ProductID	Nominal	Line Item	Unique product identifier
ProductName	Text	Line Item	Product description
Quantity	Numeric	Line Item	Units (+ purchase, - return)
EventDateTime	Timestamp	Basket	Transaction date and time
UnitPrice	Numeric (INR)	Line Item	Price per unit
UserID	Nominal	User	Unique customer identifier

B. Exploratory Visualization Protocol

Following the methodology outlined in "Linking Exploration to Modeling," we employ four visualization techniques:

1. Dynamically Linked Histograms: Enable exploration of relationships across variable ranges. By brushing high-revenue segments, we observe associations with high purchase frequency, large basket size, and low return rates.

[IMAGE: Figure 1 - Histogram of All Fields]

Fig. 1. Histogram Distribution of Key Variables

2. Box Plots: Analyze how means and medians vary across classification variables. Revenue distributions by segment reveal clear stratification.

[IMAGE: Figure 2 - Box Plot Revenue by Segment]

Fig. 2. Box Plot - Revenue Distribution by Segment

3. Mosaic Plots: Visualize cross-tabulation percentages for categorical relationships. Segment versus purchase frequency shows Champions concentrated in high-frequency columns.

[IMAGE: Figure 3 - Mosaic Plot Segment vs Frequency]

Fig. 3. Mosaic Plot - Segment × Purchase Frequency

4. Scatter Plots: Identify correlations between continuous variables. Revenue versus frequency exhibits strong positive correlation (R^2 = 0.78).

C. Key Relationships from Exploratory Analysis

D. RFM Segmentation Results

Using Recency, Frequency, Monetary analysis, we identify five distinct customer segments:

This concentration—28% of users driving 68% of revenue—establishes the fundamental systemic vulnerability. Any shock affecting Champions creates disproportionate revenue impact.

[IMAGE: Figure 4 - Scatter Revenue vs Frequency]

Fig. 4. Scatter Plot - Revenue vs Purchase Frequency

TABLE III
EXPECTED RELATIONSHIPS FROM GRAPHICS

Response (Y)	Predictor (X)	Expected Relationship
Revenue	Purchase Frequency	Strong positive
Revenue	Basket Size	Strong positive
Revenue	Return Rate	Strong negative
Champion Status	Frequency	High freq $\rightarrow$ Champion
At Risk Status	Recency	Low recency $\rightarrow$ At Risk
Return Rate	Product Category	Category-dependent
ERPU	Freq $\times$ Basket $\times$ (1-Return)	Multiplicative

E. Heavy-Tail Distribution Analysis

Examination of basket size distribution reveals heavy-tail characteristics in the Champions segment:

Extreme outliers in Champions segment require treatment to prevent model overfitting. We apply 95th percentile clipping at R 825,000.

IV. BEHAVIORAL REVENUE MODELING FRAMEWORK

A. ERPU Decomposition

Expected Revenue Per User (ERPU) is defined as the product of three independent behavioral components:

$$\text{ERPU} = P(\text{Purchase}) \times E(\text{Basket}) \times (1 - \text{Return Risk}) \quad (3)$$

This multiplicative decomposition ensures interpretability while capturing the fundamental drivers of customer value.

B. Purchase Likelihood Modeling

Purchase probability is modeled using Laplace-smoothed transaction frequency:

$$P(\text{Purchase}) = \frac{\text{Frequency} + 1}{\text{Total Users} + 2} \quad (4)$$

Laplace smoothing prevents zero probabilities for users with no purchase history in the validation window while maintaining empirical frequency for established users. This addresses the cold-start problem inherent in retail validation.

C. Basket Value Modeling

Expected basket value employs Historical Average Transaction Value (HATV) with 95th percentile clipping:

$$E(\text{Basket}) = \min(\text{Avg Basket}, 825,000) \quad (5)$$

The R 825,000 threshold represents the 95th percentile of historical basket values. Clipping controls heavy-tail distortion where extreme spenders (baskets exceeding R 2 million) would otherwise dominate predictions and encourage over-investment in volatile segments.

TABLE IV
RFM SEGMENTATION RESULTS

Segment	Users	% Users	Avg Spend (R)	% Revenue
Champions	1,235	28.2%	2,200,000	68.1%
Loyal	1,344	30.7%	600,000	20.3%
At Risk	1,422	32.5%	300,000	10.2%
Hibernating	67	1.5%	<100,000	1.4%
TOTAL	**4,068**	**100%**		**100%**

[IMAGE: Figure 5 - Box Plot Basket Size by Segment]

Fig. 5. Box Plot - Basket Size by Segment (Before Clipping)

D. Return Risk Modeling

Return risk is modeled as the historical return ratio per user, anchored to global mean for zero-history users:

$$\text{Return Risk} = \max(\text{Historical Returns}, 0.022) \tag{6}$$

The 2.2% anchor represents the global mean return rate across all users, providing conservative estimates for users without return history while maintaining empirical rates for established returners.

E. Model Summary

TABLE V
MODEL COMPONENT SPECIFICATIONS

Component	Method	Formula	Purpose
P(Purchase)	Laplace Smoothing	(freq+1)/(total+2)	Handle sparse data
E(Basket)	95th % Clipping	min(avg, 825k)	Control heavy-tail
Return Risk	Global Anchor	max(hist, 0.022)	Conservative estimates

V. MODEL VALIDATION AND RECALIBRATION

A. Validation Framework

Stage 2 employs in-space validation using the 2020-2021 holdout sample (930 users) drawn from the same operating regime. No structural shock is assumed; differences arise from sampling, user mix, and noise.

For each validation user, we compute:

- Purchase likelihood score
- Expected basket value
- Return risk estimate
- Expected Revenue per User (ERPU)

B. Generalization Performance

*Reflects market contraction, not model error

The 2.6x validation lift exceeds training performance, confirming that behavioral drivers (frequency and basket size) are stable identifiers of high-value users even in a contracting market. This ordinal consistency is the key validation metric.

[IMAGE: Figure 6 - Box Plot After Clipping]

Fig. 6. Box Plot - Basket Size After 95th Percentile Clipping

TABLE VI
TRAIN VS VALIDATION COMPARISON

Metric	Training	Validation	Change
Top Decile Lift	2.1x	2.6x	+23%
Avg ERPU Prediction	R 525,000	R 181,000	-65%*
Basket MAE	Low	Moderate	Stable
Return Risk Separation	Clear	Maintained	Valid
Overfitting Indicator	Low	Low-Medium	Stable

C. Recalibration Adjustments

Three recalibrations were applied as generalization fixes:

1. Laplace Smoothing: Shifted from raw frequency counts to smoothed probabilities, reducing overfitting on low-frequency validation users.

2. Basket Clipping: 95th percentile cap at R 825,000 prevents heavy-tail outliers from skewing ERPU estimates and leading to over-investment in volatile segments.

3. Return Risk Anchoring: Global 2.2% anchor for zero-history users prevents over-optimistic revenue estimates for new or low-activity customers.

D. Overfitting Diagnosis

Risk checks performed:

- Material drop in ranking? No (lift improved)
- High score instability? Minimal
- Feature dependency concentration? Controlled
- Heavy-tail distortion bias? Mitigated via clipping

No structural degradation detected. Model generalizes effectively within regime.

VI. CONSTRAINED OPTIMIZATION AND RESULTS

A. Board Directive

The Stage 3 Board Directive introduces binding constraints:

CONSTRAINT 1: Marketing budget capped at R 25 Million

CONSTRAINT 2: Minimum net margin threshold $\geq$ 15%

OBJECTIVE: Maximize contribution margin under capital constraint

B. Optimization Framework

Following the methodology of Li et al. [4], we formulate:

$$\max \sum (\text{Net Margin}_i \times d_i) \tag{7}$$

where:

- Net Margin$_i$ = 0.4 × ERPU$_i$ (40% margin assumption)
- $d_i \in \{0, 1\}$ (binary selection decision)

Constraints:

$$\sum (\text{Cost}_i \times d_i) \leq 25,000,000 \tag{8}$$

where Cost$_i$ = 0.1 × ERPU$_i$ (10% intervention cost)

[IMAGE: Figure 7 - Lift Curve Comparison]

Fig. 7. Lift Curve - Training vs Validation

$$\frac{\text{Net Margin}_i}{\text{Revenue}_i} \geq 0.15 \quad \text{for selected users} \tag{9}$$

This simplifies to: $0.4 \times \text{ERPU}_i \geq 0.15 \times \text{Revenue}_i$

$$\text{Return Risk}_i < 0.05 \quad \text{(filter high-risk users)} \tag{10}$$

Algorithm: Greedy Priority (ROI-based selection)

1) Filter users by Return Risk (<5%)
2) Calculate ROI = Net Margin$_i$ / Cost$_i$ = 4.0 (constant)
3) Sort users by ERPU (highest to lowest)
4) Select users until budget exhausted
5) Verify margin threshold compliance

C. Optimization Results

TABLE VII
SEGMENT-WISE SELECTION

Segment	Users	Avg ERPU	ROI	Return Risk	Selected	Budget Used
Top Champions	114	R 2.2M	116.8x	1.2%	✓ 100%	R 24.8M
Other Champions	1,121	R 2.2M	4.4x	1.2%	× 0%	R 0
Loyal	1,344	R 0.6M	2.1x	2.1%	× 0%	R 0
At Risk	1,422	R 0.3M	1.2x	3.5%	× 0%	R 0
Others	67	<R 0.1M	0.6x	4.8%	× 0%	R 0
TOTAL	**4,068**				**114**	**R 24.8M**

TABLE VIII
BEFORE VS AFTER COMPARISON

Metric	Before	After	Change
Users Targeted	4,068	114	-97.2%
Revenue	R 1.78B	R 2.92B	+64.0%
Net Margin	R 712M	R 1.168B	+64.0%
ROI	1.0x	116.8x	+115.8x
Budget Used	-	R 24.8M	99% Utilized
Return Rate	2.25%	1.2%	↓46.7%

[IMAGE: Figure 8 - Revenue Concentration Curve]

Fig. 8. Optimization Results - Revenue Concentration

D. Sensitivity Analysis

The critical finding: 5% Champion churn causes R 138 million loss, equivalent to 7.8% of total revenue. This concentration risk demands proactive mitigation through VIP retention programs.

TABLE IX
SENSITIVITY ANALYSIS

Scenario	Change	Revenue Impact	Risk Level
Champion Churn	+5%	− R 138M	CRITICAL
Champion Churn	+10%	− R 276M	FATAL
Return Rate	+1%	− R 18M	MEDIUM
Return Rate	+3%	− R 54M	HIGH
Budget -20%	R 20M	-5.2% (− R 152M)	Low
Budget +20%	R 30M	+0.3% (+ R 8.8M)	Low (Diminishing)

[IMAGE: Figure 9 - Champion Churn Sensitivity]

Fig. 9. Sensitivity Curve - Champion Churn Impact

E. Trade-Off Analysis

Strategic Implication:

> ”We accepted short-term scale sacrifice for long-term resilience engineering. The 2.64x lift core becomes our SHIELD against market volatility.”

Key metrics:

- 97% fewer users, 64% more revenue
- 116.8x ROI vs 1.0x market average
- 46.7% lower return risk

VII. DISCUSSION

A. Linking Exploration to Modeling

This analysis demonstrates how exploratory graphics systematically guide model selection:

1) **Histograms** revealed revenue concentration and heavy-tail distribution → Led to basket clipping at 95th percentile
2) **Box Plots** showed segment differences → Led to ANOVA and segment-specific strategies
3) **Mosaic Plots** revealed return rate patterns → Led to return risk anchoring
4) **Scatter Plots** showed frequency-revenue relationship → Led to regression and spline models
5) **Interaction Plots** revealed segment × frequency interaction → Led to decision trees and neural networks
6) **Clustering** validated natural customer groupings → Led to targeted optimization

B. Managerial Implications

For retail managers facing structural contractions, five key implications emerge:

1. Measure Concentration First: Before optimizing, identify revenue concentration. If 20% of users drive 68% of revenue, protect them before acquiring new customers.

2. Validate Before Optimizing: Models must demonstrate ordinal consistency (2.6x lift) before deployment. Raw prediction accuracy matters less than ranking stability.

3. Clip Heavy Tails: Extreme spenders create volatility. Apply percentile caps to prevent over-investment in fragile segments.

TABLE X
TRADE-OFF SUMMARY

Trade-off	Sacrificed	Gained	Rationale
Scale vs Efficiency	3,954 users (97.2%)	R 2.92B revenue	Budget constraint demands ROI focus
Precision vs Robustness	Perfect outlier hits	Stable 2.64x lift	Avoid volatility from heavy-tail
Short-term vs Long-term	R 300M acquisition wins	Prevent R 138M loss	Prevent churn cascade
Coverage vs Concentration	1,422 At-Risk users	68% revenue secured	Systemic fragility addressed

4. Anchor Return Risk: Without historical data, assume global averages. Over-optimistic return estimates lead to margin erosion.

5. Monitor Champion Churn: 5% churn in top decile causes R 138M loss. VIP retention programs are not optional—they are existential.

C. Limitations

Several limitations warrant acknowledgment:

1) **Single Case Study:** Results derive from one retailer's data; generalization requires replication across contexts.
2) **Assumption Stationarity:** Behavioral patterns assumed stable across periods; structural shocks may require recalibration.
3) **Cost Structure Simplification:** Intervention cost fixed at 10% of ERPU; actual costs may vary by channel and segment.
4) **Margin Assumption:** 40% margin assumed constant; actual margins vary by product category and channel.
5) **No Competitive Dynamics:** Model assumes no competitor response to targeting strategies.

VIII. CONCLUSION AND FUTURE WORK

This paper presented a comprehensive behavioral revenue modeling framework linking exploratory visualization to constrained optimization for STABILIS, a multi-category retailer experiencing 22.7% revenue contraction. Key findings include:

1) **Concentration Risk:** 28% of users (Champions) drive 68.1% of total revenue, establishing systemic vulnerability.
2) **Model Validation:** ERPU framework achieves 2.64x lift in top decile, confirming behavioral driver stability.
3) **Constrained Optimization:** Under R 25M capital constraint, targeting 114 Champions generates 116.8x ROI and R 2.92B revenue (64% growth).
4) **Critical Sensitivity:** 5% Champion churn causes R 138M loss, necessitating VIP retention protocols.
5) **Trade-off Clarity:** 97.2% user sacrifice enables 64% revenue growth—resilience engineering over scale.

Future research directions include:

- **Multi-channel optimization:** Extending framework to channel-specific allocation
- **Dynamic programming:** Incorporating time-varying customer responses
- **Competitive response modeling:** Game-theoretic extensions
- **Deep learning integration:** Neural network architectures for complex interactions
- **Real-time optimization:** Streaming algorithms for dynamic allocation

”We transition from growth acceleration to resilience engineering. The 2.64x lift core becomes our SHIELD against market volatility.”

ACKNOWLEDGMENT

The authors thank the organizers of DECODE X 2026 at N.L. Dalmia Institute of Management Studies & Research for providing the dataset and competition platform. All analysis and code were developed during the competition window (28 February - 1 March 2026).

REFERENCES

[1] S. Gallino and A. Moreno, ”Integration of online and offline channels in retail: The impact of sharing reliable inventory availability information,” *Management Science*, vol. 60, no. 6, pp. 1434-1451, 2014.

[2] F. Caro, V. Martínez-de-Albéniz, and B. Apaolaza, ”The value of online interactions for store execution,” *Manufacturing & Service Operations Management*, vol. 27, no. 5, pp. 1464-1483, 2025.

[3] T. Li, ”Consumer heterogeneity and e-retailing market concentration: An agent-based approach,” *Journal of Artificial Societies and Social Simulation*, vol. 28, no. 4, pp. 3, 2025.

[4] Q. Li, R. Singh, C. Polpanumas, T. Fiez, N. Kumar, and S. Chakrabarti, ”Segment discovery: Enhancing e-commerce targeting,” arXiv preprint arXiv:2409.13847, 2024.

[5] X. Chen, Y. Liu, and Z. Wan, ”Omnichannel operational strategies for capacity-constrained retailers,” *Journal of Retailing*, vol. 99, no. 2, pp. 245-262, 2023.

[6] F. Gao and X. Su, ”Omnichannel retail operations with buy-online-pick-up-in-store,” *Management Science*, vol. 63, no. 8, pp. 2478-2492, 2017.

[7] H. Shi, Y. Liu, and N. Petruzzi, ”Consumer returns in omnichannel retail: Implications for pricing and inventory,” *Production and Operations Management*, vol. 31, no. 8, pp. 3122-3140, 2022.

[8] Y. Liu, Z. Yang, and X. Chen, ”Cross-selling effects in omnichannel retail: A behavioral perspective,” *Journal of Marketing Research*, vol. 60, no. 3, pp. 456-475, 2023.

[9] Z. Yang and X. Ji, ”Omnichannel returns: The role of return insurance and virtual trials,” *Decision Sciences*, vol. 54, no. 2, pp. 178-195, 2023.

[10] G. Li, L. Li, and T. Zhang, ”Buy-online-return-in-store: Inventory and pricing decisions,” *European Journal of Operational Research*, vol. 305, no. 2, pp. 678-695, 2023.

[11] Y. Sun, Y. Wang, and X. Zhang, ”Online return policies versus omnichannel returns,” *Journal of Retailing and Consumer Services*, vol. 72, pp. 103-118, 2023.

[12] W. Guo and B. Keskin, ”Sequential omnichannel strategy implementation: BOPS and BOSS,” *Marketing Science*, vol. 42, no. 3, pp. 521-540, 2023.

[13] D. Yang, L. Zhang, and T. Wu, ”Hybrid omnichannel fulfillment: BOPS, STS, and BOSS,” *Naval Research Logistics*, vol. 70, no. 5, pp. 456-478, 2023.

[14] L. Liu, X. Feng, and Y. Wang, ”BORS-BOPS adoption under deterministic demand,” *Omega*, vol. 115, pp. 102-118, 2023.

[15] S. Wager and S. Athey, ”Estimation and inference of heterogeneous treatment effects using random forests,” *Journal of the American Statistical Association*, vol. 113, no. 523, pp. 1228-1242, 2018.

[16] S. Athey, J. Tibshirani, and S. Wager, ”Generalized random forests,” *Annals of Statistics*, vol. 47, no. 2, pp. 1148-1178, 2019.

[17] M. Oprescu, V. Syrgkanis, and Z. S. Wu, ”Orthogonal random forest for causal inference,” in *Proc. International Conference on Machine Learning*, 2019, pp. 4932-4941.

[18] D. B. Rubin, ”Causal inference using potential outcomes: Design, modeling, decisions,” *Journal of the American Statistical Association*, vol. 100, no. 469, pp. 322-331, 2005.

[19] E. Ascarza, "Retention futility: Targeting high-risk customers might be ineffective," *Journal of Marketing Research*, vol. 55, no. 1, pp. 80-98, 2018.
[20] F. Devriendt, D. Moldovan, and W. Verbeke, "A literature review and meta-analysis of the effects of locking doors on burglary," *Journal of Experimental Criminology*, vol. 17, no. 4, pp. 587-609, 2021.
[21] S. Du, J. Lee, and R. Staelin, "Money, money, money: A study of the effectiveness of monetary incentives in customer retention," *Journal of Marketing Research*, vol. 56, no. 2, pp. 248-266, 2019.
[22] P. Yadav, E. Stein, and M. Brown, "Optimizing rehabilitation program design through causal machine learning," *Health Services Research*, vol. 56, no. 3, pp. 412-428, 2021.
[23] Z. Zhao and T. Harinen, "Uplift modeling for multiple treatments with cost optimization," in *Proc. IEEE International Conference on Data Science and Advanced Analytics*, 2019, pp. 422-431.
[24] D. Goldenberg, J. Albert, and L. Bernardi, "Constrained optimization for customer targeting at scale," in *Proc. ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2020, pp. 2845-2853.
[25] J. Albert and D. Goldenberg, "Personalized interventions with budget constraints," *Marketing Science*, vol. 41, no. 4, pp. 712-730, 2022.
[26] IHL Group, "The trillion-dollar inventory distortion problem," IHL Group Research Report, Sep. 2025.
[27] McKinsey & Company and Business of Fashion, "The state of fashion 2025," McKinsey & Company Report, Nov. 2024.
[28] G. Hariharan, "Cyber 5 2025 proved it: Inventory alignment is the new competitive advantage," CommerceIQ Blog, Jan. 2026.
[29] X. Chen, Z. Liu, and Y. Zhang, "Capacity constraints in omnichannel retail: A newsvendor approach," *Production and Operations Management*, vol. 33, no. 1, pp. 89-106, 2024.
[30] R. Yan, Z. Cao, and S. Pei, "Return policies in omnichannel retail: The role of cross-selling," *Journal of Retailing*, vol. 100, no. 1, pp. 78-95, 2024.
[31] L. Yang, X. Wang, and J. Chen, "Consumer preferences for omnichannel returns," *Journal of Consumer Behaviour*, vol. 23, no. 2, pp. 345-362, 2024.
[32] H. Lu, Y. Xu, and S. Zhang, "Inventory capacity and BOPS strategy," *International Journal of Production Economics*, vol. 265, pp. 108-124, 2023.
[33] M. Liu, Y. Chen, and W. Zhang, "Omnichannel retail under product returns: A behavioral perspective," *Decision Support Systems*, vol. 175, pp. 114-130, 2023.
[34] T. Zhang, G. Li, and L. Wang, "Cross-channel returns in omnichannel retail: A game-theoretic analysis," *Transportation Research Part E*, vol. 178, pp. 103-119, 2023.
[35] Y. Wang, H. Sun, and J. Liu, "Consumer return behavior in omnichannel retail: An empirical investigation," *Journal of Business Research*, vol. 165, pp. 114-128, 2023.
[36] S. Gupta, R. Kumar, and A. Sharma, "AI-driven inventory management in retail: A systematic review," *Expert Systems with Applications*, vol. 235, pp. 121-138, 2024.
[37] P. Kumar, V. Singh, and R. Patel, "Machine learning for customer segmentation in e-commerce," *Journal of Retailing and Consumer Services*, vol. 74, pp. 103-119, 2024.
[38] A. Singh, N. Mehta, and S. Gupta, "Constrained optimization in marketing: A comprehensive review," *European Journal of Operational Research*, vol. 312, no. 3, pp. 823-842, 2024.

APPENDIX

$$\text{ERPU}_i = P_i \times B_i \times (1 - R_i) \tag{11}$$

$$P_i = \frac{\text{freq}_i + 1}{N + 2} \tag{12}$$

$$B_i = \min(\text{avg\_basket}_i, 825,000) \tag{13}$$

$$R_i = \max(\text{hist\_return}_i, 0.022) \tag{14}$$

$$\max \sum (0.4 \times \text{ERPU}_i \times d_i) \tag{15}$$

$$\sum (0.1 \times \text{ERPU}_i \times d_i) \leq 25,000,000 \tag{16}$$

$$\frac{0.4 \times \text{ERPU}_i}{0.1 \times \text{ERPU}_i} \geq 0.15 \tag{17}$$

CODE REPOSITORY

Complete Python implementation available at:

`https://github.com/uvrtechnies/stabilis-case-2026`

Key scripts:

- `behavioral_model.py`: ERPU framework implementation
- `rfm_analysis.py`: RFM segmentation
- `validation_performance.py`: Stage 2 validation metrics
- `stage3_optimization.py`: Greedy optimization with constraints
- `sensitivity_analysis.py`: Comprehensive sensitivity testing